

MICROSOFT MOVIES ANALYSIS

- Student name: Frida Oyucho
- Student pace: full time
- Instructor name: William Okombo/Noah Kandie

Data Preparation

Steps taken

- Importing of merged data
- Dropped unnecessary columns
- Removed special characters and whitespace from Production_budget & worldwide_gross columns
- Changed Data Types for Production_budget & worldwide_gross columns
- Dropped rows in runtime_minutes column
- Replaced Null values with MISSING for Movies column
- Removed the missing values which were denoted by zero for Production_budget & worldwide_gross columns
- Identified and removed duplicates
- Removed extreme outliers using IQR

```
In [2]: # Importing necessary Libraries  
import pandas as pd  
import numpy as np  
import matplotlib.pyplot as plt
```

In [3]: *# Importing them merged dataset and converting it to a dataframe*

```
df = pd.read_csv('FinalMergedData.csv', index_col=0)
df.head()
```

Out[3]:

	genre_ids	id_x	original_language	original_title	popularity	release_date_x	title
0	[12, 14, 10751]	12444	en	Harry Potter and the Deathly Hallows: Part 1	33.533	2010-11-19	Harry Potter and the Deathly Hallows: Part 1
1	[14, 12, 16, 10751]	10191	en	How to Train Your Dragon	28.734	2010-03-26	How to Train Your Dragon
2	[12, 28, 878]	10138	en	Iron Man 2	28.515	2010-05-07	Iron Man 2
3	[28, 878, 12]	27205	en	Inception	27.920	2010-07-16	Inception
4	[12, 14, 10751]	32657	en	Percy Jackson & the Olympians: The Lightning T...	26.691	2010-02-11	Percy Jackson & the Olympians: The Lightning T...

In [4]: *# Preview column names*

```
df.columns
```

Out[4]: Index(['genre_ids', 'id_x', 'original_language', 'original_title', 'popularity', 'release_date_x', 'title', 'vote_average', 'vote_count', 'tconst', 'primary_title', 'start_year', 'runtime_minutes', 'genres', 'id_y', 'release_date_y', 'movie', 'production_budget', 'domestic_gross', 'worldwide_gross'], dtype='object')

In [5]: *# Dropping unnecessary Columns*

```
df.drop(['genre_ids', 'id_x', 'original_title', 'release_date_x', 'id_y', 'release_date_y'], axis=1, inplace=True)
df.shape
```

Out[5]: (21078, 12)

```
In [7]: # Checking information about the dataset
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 21078 entries, 0 to 21077
Data columns (total 12 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   original_language     21078 non-null  object
 1   popularity             21078 non-null  float64
 2   title                  21078 non-null  object
 3   vote_average          21078 non-null  float64
 4   tconst                 21078 non-null  object
 5   primary_title         21078 non-null  object
 6   start_year            21078 non-null  int64
 7   runtime_minutes       19452 non-null  float64
 8   genres                 20779 non-null  object
 9   movie                  5782 non-null   object
10   production_budget     5782 non-null   object
11   worldwide_gross       5782 non-null   object
dtypes: float64(3), int64(1), object(8)
memory usage: 2.1+ MB
```

```
In [8]: # Removing $ sign from Production_budget & worldwide_gross
df['production_budget']=df['production_budget'].str.replace('$', '')
df['production_budget']

df['worldwide_gross']=df['worldwide_gross'].str.replace('$', '')
```

```
In [9]: # Removing ',' sign from Production_budget & worldwide_gross
df['production_budget']=df['production_budget'].str.replace(',', '')
df['production_budget']

df['worldwide_gross']=df['worldwide_gross'].str.replace(',', '')
```

```
In [10]: df['production_budget'].dtype
```

```
Out[10]: dtype('O')
```

```
In [11]: # Changing objects to int for Production Budget & Worldwide_gross
# Replace NaN with 0
df['production_budget'] = df['production_budget'].fillna(0)
df['production_budget'] = df['production_budget'].astype(int)

# Changing objects to int for Production Budget & Worldwide_gross
# Replace NaN with 0
df['worldwide_gross'] = df['worldwide_gross'].fillna(0)
df['worldwide_gross'] = df['worldwide_gross'].astype(float)
```

```
In [12]: # Handling missing Values
# Checking the proportion of missing data
df.isnull().mean()
```

```
Out[12]: original_language    0.000000
popularity                  0.000000
title                      0.000000
vote_average               0.000000
tconst                     0.000000
primary_title              0.000000
start_year                 0.000000
runtime_minutes            0.077142
genres                     0.014185
movie                      0.725686
production_budget          0.000000
worldwide_gross            0.000000
dtype: float64
```

```
In [13]: # Dropping rows in runtime_minutes column because of the low proportion of
df = df.dropna(subset=['runtime_minutes', 'genres'])
# df.isnull().mean()
```

```
In [14]: # Replacing Null values with the word MISSING for Movies
df['movie']=df['movie'].fillna('missing')
df.isnull().mean()
```

```
Out[14]: original_language    0.0
popularity                  0.0
title                      0.0
vote_average               0.0
tconst                     0.0
primary_title              0.0
start_year                 0.0
runtime_minutes            0.0
genres                     0.0
movie                      0.0
production_budget          0.0
worldwide_gross            0.0
dtype: float64
```

```
In [15]: df.shape
```

```
Out[15]: (19329, 12)
```

```
In [16]: #Removing the missing values which are denoted by zero
df=df[(df['production_budget']!=0)]
df=df[(df['worldwide_gross']!=0)]
df.shape
```

```
Out[16]: (4853, 12)
```

```
In [17]: # Handling duplicate Records
# Checking for duplicates in the dataset
df.duplicated().any()

# Dropping the duplicates
df.drop_duplicates(inplace=True, keep = 'first')
```

```
In [18]: df.shape
```

```
Out[18]: (4853, 12)
```

```
In [19]: # Handling Outliers
df.describe()
```

```
Out[19]:
```


	popularity	vote_average	start_year	runtime_minutes	production_budget	worldwi
count	4853.000000	4853.000000	4853.000000	4853.000000	4.853000e+03	4.85
mean	4.014960	5.696023	2012.128992	91.926437	3.338989e+07	9.89
std	4.740305	1.570870	2.454837	23.707368	4.328410e+07	1.83
min	0.600000	0.000000	2010.000000	3.000000	5.000000e+03	2.60
25%	0.840000	5.000000	2010.000000	83.000000	6.000000e+06	6.60
50%	1.985000	5.800000	2011.000000	91.000000	1.800000e+07	3.29
75%	6.136000	6.600000	2013.000000	101.000000	4.000000e+07	1.04
max	50.289000	10.000000	2020.000000	495.000000	4.250000e+08	2.77

```
In [20]: import seaborn as sns
```

```
In [23]: # Visualizing Outliers
sns.boxplot(df['popularity']);
```

C:\Users\Fridah.Oyucho\AppData\Local\anaconda3\envs\learn-env\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
warnings.warn(
```



```
In [24]: # Position of the Outlier values
x = np.where(df['runtime_minutes'] > 150)
print(x)
```


```
(array([ 261,  322,  391,  409,  473,  474,  478,  479,  604,  876,  919,
        953,  957,  961, 1515, 1520, 1522, 1700, 1796, 1839, 1909, 1914,
        1919, 1924, 1929, 1979, 2009, 2074, 2075, 2096, 2098, 2107, 2271,
        2364, 2492, 2605, 2612, 2683, 2687, 2786, 2828, 2832, 2834, 2939,
        2943, 2946, 2950, 3007, 3129, 3390, 3408, 3492, 3506, 3510, 3568,
        3844, 3849, 3853, 3882, 3946, 4135, 4361, 4392, 4764, 4770, 4775,
        4835], dtype=int64),)
```

```
In [21]: sns.boxplot(df['runtime_minutes'])
```

C:\Users\Fridah.Oyucho\AppData\Local\anaconda3\envs\learn-env\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
warnings.warn(
```

```
Out[21]: <AxesSubplot:xlabel='runtime_minutes'>
```



```
In [25]: #Creating a function to remove the outliers
def remove_outliers_usingIQR(df,column):
    q1 = df[column].quantile(0.25)
    q3 = df[column].quantile(0.75)
    iqr = q3 - q1
    lower_bound = q1 - 1.5*iqr
    upper_bound = q3 + 1.5*iqr
    return df[(df[column]>= lower_bound) & (df[column] <= upper_bound)]
```

```
In [26]: #removing outliers using the created function
df = remove_outliers_usingIQR(df,'runtime_minutes')
df = remove_outliers_usingIQR(df,'production_budget')
df = remove_outliers_usingIQR(df,'worldwide_gross')
```

```
In [27]: df = remove_outliers_usingIQR(df,'popularity')
```

```
In [28]: df.describe()
```

```
Out[28]:
```

	popularity	vote_average	start_year	runtime_minutes	production_budget	worldwi
count	3454.000000	3454.000000	3454.000000	3454.000000	3.454000e+03	3.45
mean	2.803139	5.559873	2012.102490	91.269543	2.002213e+07	3.89
std	2.749238	1.621166	2.355448	12.667780	1.990075e+07	4.55
min	0.600000	0.000000	2010.000000	56.000000	5.000000e+03	2.60
25%	0.665250	4.800000	2011.000000	84.000000	5.000000e+06	4.02
50%	1.552000	5.600000	2011.000000	90.000000	1.400000e+07	2.01
75%	3.836000	6.600000	2012.000000	99.000000	3.000000e+07	5.89
max	11.571000	10.000000	2020.000000	128.000000	9.000000e+07	1.87

```
In [29]: df.shape
```


```
Out[29]: (3454, 12)
```

```
In [30]: # Visualizing boxplots after removing Outliers
sns.boxplot(df['production_budget'])
```

C:\Users\Fridah.Oyucho\AppData\Local\anaconda3\envs\learn-env\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
warnings.warn(
```

```
Out[30]: <AxesSubplot:xlabel='production_budget'>
```





```
In [28]: # Visualizing boxplots after removing Outliers
sns.boxplot(df['runtime_minutes'])
```

C:\Users\Fridah.Oyucho\AppData\Local\anaconda3\envs\learn-env\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
warnings.warn(
```

Out[28]: <AxesSubplot:xlabel='runtime_minutes'>




```
In [31]: # Visualizing boxplots after removing Outliers
sns.boxplot(df['worldwide_gross'])
```

C:\Users\Fridah.Oyucho\AppData\Local\anaconda3\envs\learn-env\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
warnings.warn(
```


Out[31]: <AxesSubplot:xlabel='worldwide_gross'>



```
In [32]: # Visualizing boxplots after removing Outliers
sns.boxplot(df['popularity']);
```

C:\Users\Fridah.Oyucho\AppData\Local\anaconda3\envs\learn-env\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
warnings.warn(
```



```
In [33]: df.describe()
```

Out[33]:

	popularity	vote_average	start_year	runtime_minutes	production_budget	worldwi
count	3454.000000	3454.000000	3454.000000	3454.000000	3.454000e+03	3.45
mean	2.803139	5.559873	2012.102490	91.269543	2.002213e+07	3.89
std	2.749238	1.621166	2.355448	12.667780	1.990075e+07	4.55
min	0.600000	0.000000	2010.000000	56.000000	5.000000e+03	2.60
25%	0.665250	4.800000	2011.000000	84.000000	5.000000e+06	4.02
50%	1.552000	5.600000	2011.000000	90.000000	1.400000e+07	2.01
75%	3.836000	6.600000	2012.000000	99.000000	3.000000e+07	5.89
max	11.571000	10.000000	2020.000000	128.000000	9.000000e+07	1.87

```
In [34]: df.shape
```

Out[34]: (3454, 12)

```
In [35]: # Exporting Output to CSV file
df.to_csv('CleanedData.csv', index=False)
```

Conclusion on Handling Outliers

We managed to remove extreme but not all outliers

